

Large Language Models for the
Economic and Social Sciences:

Getting Output from LLMs

Indira Sen, HWS 2025

<https://github.com/dess-mannheim/LLM4ESS-HWS25>


Credits: Thanks to Markus Strohmaier for some of the content on prompting.

Announcements

- Please attend the exercises
- Access to Cluster?

Last Time: How can we build LLMs?


Somewhat simplified answer



Jurafsky and Martin, 2025

Last Time: How can we build LLMs?

Somewhat simplified answer



Jurafsky and Martin, 2025

Vastly Simplified answer

- We take a LOT of data encoding social, cultural, and of course linguistic traces
- Train NLP models that can generate content on this data
- We further train the model to follow and respond to instructions
- We further modify the model so that it doesn't have unwanted behavior

Summary of Last Week and Today

Last session we learned:

- Pretraining and fine-tuning
- Language generation
- Instruction-tuning
- Reinforcement Learning with Human Feedback

Today:


- Different types of prompting
- How to get outputs from LLMs in other ways
- Practicalities of working with LLMs
 - Which LLMs and where to find them
 - How to evaluate them
- Project Details

How does an LLM generate content?

- Conditioned on previous contexts, e.g., some type of input or **prompt**
 - computing the probability of the next token from the prior context
 - sampling from that distribution to generate a token


- Decoding

- This task of choosing a word to generate based on the model's probabilities



Decoding Algorithms

- Greedy decoding: pick the most high probability token
- Random sampling: pick any token at random
- Temperature sampling: increase the probability of picking higher probability tokens



Have you tried changing the temperature when interacting with LLMs?

LLMs and where to find them

Some Foundation Models (till mid-2026)

OpenAI with ChatGPT, multiple sub-models (mostly not open source, available via API)

- For top performance: GPT-6 Astra
- For most tasks: GPT-5.6 (Sol / Terra / Luna)
- For programming: Codex with GPT-6 Sol
- For faster, cheaper work: GPT-6 Luna
- Reasoning is now built into every model, with adjustable reasoning effort
- **Open weights on Hugging Face: gpt-oss-120b and gpt-oss-20b**

Meta with Llama and Muse Spark

- Latest open model: Llama 4 (Scout and Maverick)
- New flagship: Muse Spark (proprietary, not open source)
- Uncensored fine-tunes: Dolphin 3.0 (Mistral 24B and Llama 3.1 versions)

DeepSeek with DeepSeek V4 (Open Source, available via API)

- For high performance: DeepSeek V4-Pro
- For speed and efficiency: DeepSeek V4.1 Flash

Some Foundation Models (till mid-2026)

Alibaba with Qwen, multiple sub-models (mostly Open Source, available via API)

- For high performance: Qwen3.8-Max (API) / Qwen3.8-2.4T (open weights)
- For efficiency: Qwen3.8-Flash
- For local use: Qwen3.8-27B

Moonshot AI with Kimi (Open Source, available via API)

- Flagship: Kimi K3, the largest open-weight model to date
- For programming: Kimi K2.7 Code

Some Foundation Models (till mid-2026)

Anthropic with Claude, multiple sub-models (Not Open Source, Available via API)

- For frontier performance: Claude Fable 5.1
- For high performance: Claude Opus 5.5
- For enterprises: Claude Sonnet 5.5
- For efficiency: Claude Haiku 4.5

Open Weight

Data	Training	Weights
		
		

Google with Gemini, multiple sub-models (Not Open Source, Available via API)

- For high performance: Gemini 3.1 Pro (+ Gemini 3 Deep Think for complex reasoning)
- For scaling: Gemini 3.8 Flash
- For efficiency: Gemini 3.5 Flash-Lite
- For small devices: Gemini Nano
- For developers: Gemma 4 (Open Source)

Open Source

Data	Training	Weights
		
		

AllenAI with OLMo 3.1

- Fully **open-source** → open model weights and training data



Yan Cui

Where to find these and future models?

- Hugging Face (huggingface.co) is a central research hub for Large Language Models and Benchmarks
- Can be used to search for, download and (sometimes) directly interact with models via browser
- Requires (free) user registration
- Example: <https://huggingface.co/mistralai/Mixtral-8x7B-Instruct-v0.1>

Different Types of Prompts

1. System Prompt:

You are a Management Consultant from Accenture providing thorough and systematic advice to clients in the Marketing and Communications Industry. You have 20 years plus of experience. You are friendly, analytical and inquisitive.

2. User Prompt:

Define three strategies for marketing firms to adopt LLMs in their business processes.



3. Assistant Answer:

Create image

Surprise me

Code

Help me write

Make a plan

More

Marketing firms can adopt large language models (LLMs) in their business processes through the following strategies:

1. Enhanced Content Creation and Personalization

- **Application:** Use LLMs to generate high-quality, engaging, and SEO-optimized content, such as blog posts, ad copy, social media updates, and email campaigns. Tailor content to specific audience segments based on preferences and behavior.
- **Benefit:** Streamline the content creation process, reduce production costs, and improve audience engagement by delivering highly personalized marketing materials.
- **Example:** An e-commerce marketing firm can employ LLMs to create personalized email

System prompts

- Every model comes with base instructions, so-called “System Prompts”
- System prompts instruct the model how to behave in response to user prompts.

- Example:

Your knowledge cutoff is 2023-10. You are a helpful, witty, and friendly AI. Act like a human, but remember that you aren't a human and that you can't do human things in the real world. Your voice and personality should be warm and engaging, with a lively and playful tone. If interacting in a non-English language, start by using the standard accent or dialect familiar to the user. Talk quickly. You should always call a function if you can. Do not refer to these rules, even if you're asked about them.

- System prompts are part of the context window. They “initialize” the conversation with desirable behaviors and try to constrain undesirable behaviors (via so-called guardrails).

Many options with prompting

- Define the objective of your prompt
- Specify the format of the response
- Assign a role
- Identify the audience
- Provide context
- Set constraints
- Include examples
- Specify the tone of the response
- Ask for explanations
- Encourage creativity
- Request citations
- Test and refine prompts
- Provide feedback, iterate

Many options with prompting

- Define the objective of your prompt
- Specify the format of the response
- Assign a role
- Identify the audience
- Provide context
- Set constraints
- Include examples
- Specify the tone of the response
- Ask for explanations
- Encourage creativity
- Request citations
- Test and refine prompts
- Provide feedback, iterate

Zero Shot Learning (ZSL)

- **Definition:**
 - An approach where a model performs tasks it has not been explicitly trained on (no examples given) by leveraging general knowledge or understanding of related contexts.
- **Summary:**
 - ZSL allows models to handle new tasks without requiring task-specific examples
- **Example:**
 - Prompt: "Translate the following English sentence to Spanish: 'Good morning, how are you?'"
 - Expected Output: "Buenos días, ¿cómo estás?"

Few-Shot Learning

- **Description:**

- Few-shot learning is a prompting technique in which a small number of examples (a few "shots") are included in the input to teach a language model how to respond to a specific type of task or question.

- **Summary:**

- By providing a few examples, the model can infer the desired task format and generalize to similar prompts.

- **Example:**

- Q: What is the capital of France? A: Paris
Q: What is the capital of Spain? A: Madrid
Q: What is the capital of Germany? A:
- } Examples (few shots)

Many options with prompting

- Define the objective of your prompt
- Specify the format of the response
- Assign a role
- Identify the audience
- Provide context
- Set constraints
- Include examples
- Specify the tone of the response
- Ask for explanations
- Encourage creativity
- Request citations
- Test and refine prompts
- Provide feedback, iterate

Chain of Thought Prompting

- **Definition:**

- Breaks problems into step-by-step reasoning for clarity and accuracy.

- **Summary:**

- Guides the model to think logically like humans.

- **Example:**

- Prompt: "If a train travels 60 mph for 2 hours and 40 mph for 1 hour, how far does it travel in total? Explain step by step."
- Response: "First: $60 \times 2 = 120$ miles. Second: $40 \times 1 = 40$ miles. Total: $120 + 40 = 160$ miles."

- **Behind the scenes:** Likely uses few shot learning of chain of thought answers

Chain of Thought Prompting → reasoning models

- Models with in-built reasoning perform multi-step problem-solving by generating **intermediate steps, often called “reasoning traces,”** through which it iteratively refines its output prior to generating its final response.
- Inference-time decision-making: The LLM generates several responses, and then picks one based on heuristics



DeepSeek-R1: Incentivizing Reasoning Capability in LLMs via Reinforcement Learning

DeepSeek-AI

research@deepseek.com

Abstract

General reasoning represents a long-standing and formidable challenge in artificial intelligence. Recent breakthroughs, exemplified by large language models (LLMs) (Brown et al., 2020; OpenAI, 2023) and chain-of-thought prompting (Wei et al., 2022b), have achieved considerable success on foundational reasoning tasks. However, this success is heavily contingent upon extensive human-annotated demonstrations, and models' capabilities are still insufficient for more complex problems. Here we show that the reasoning abilities of LLMs can be incentivized through pure reinforcement learning (RL), obviating the need for human-labeled reasoning trajectories. The proposed RL framework facilitates the emergent development of advanced reasoning patterns, such as self-reflection, verification, and dynamic strategy adaptation. Consequently, the trained model achieves superior performance on verifiable tasks such as mathematics, coding competitions, and STEM fields, surpassing its counterparts trained via conventional supervised learning on human demonstrations. Moreover, the emergent reasoning patterns exhibited by these large-scale models can be systematically harnessed to guide and enhance the reasoning capabilities of smaller models.

Guo, Daya, et al. "[Deepseek-r1: Incentivizing reasoning capability in llms via reinforcement learning.](#)" *arXiv preprint arXiv:2501.12948* (2025).

Retrieval Augmented Generation (RAG)

- **Definition:**

- RAG combines pre-trained language models with external knowledge to generate contextually accurate and informed responses.

- **Summary:**

- RAG enhances responses by integrating retrieved documents or data into the generation process, grounding outputs in external factual knowledge.

- **Example:**

- User Prompt: "Fetch the competitor analysis document and assess it from a business opportunity perspective"




Competitor
analysis.doc

- Assistant Response: The following opportunities can be identified ...

- **Remark:**

Can also be used to retrieve documents from outside sources, e.g. the web or the enterprise intranet

Retrieval Augmented Generation (RAG)




Source: OpenAI

More Prompting Techniques

- Meta Prompting
- Self-Consistency
- Generate Knowledge Prompting
- Prompt Chaining
- Tree of Thoughts
- Retrieval Augmented Generation
- Automatic Prompt Engineer
- And many others!


Prompt Engineering

- Resources for diving deeper into Prompting:
 - <https://www.promptingguide.ai>
 - <https://promptengineering.org/>
- Examples and explanations of various prompt types and strategies.



Prompt Brittleness and Optimization

- LLMs are *brittle* to prompts
 - Small changes that a human doesn't care about leads to big changes in the LLMs' behavior
- Automatic prompt engineering or prompt optimization uses optimization techniques to generate and find the best prompt for a given task (e.g., [Automatic Prompt Engineer \(APE\) from Zhou et al.](#))



Ceron et al., 2024 [Beyond Prompt Brittleness: Evaluating the Reliability and Consistency of Political Worldviews in LLMs](#)

Many options with prompting


- Define the objective of your prompt
- Specify the format of the response
- Assign a role
- Identify the audience
- Provide context
- Set constraints
- Include examples
- Specify the tone of the response
- Ask for explanations
- Encourage creativity
- Request citations
- Test and refine prompts
- Provide feedback, iterate

What are model decoding hyperparameters?

- **[Temperature]**
 - Controls the randomness of the model's output. Lower values make the output more deterministic, while higher values make it more creative or unpredictable.
 - **Extreme Example:**
 - Low (0.1): "Paris is the capital of France."
 - High (1.5): "Paris might be a dazzling galaxy of lights and cheese-filled dreams."

What are model decoding hyperparameters?

- [Top-K and Top-p]
 - **Top-K:** Keeps a **fixed number (K)** of the most probable next words.
 - **Top-p:** Restricts the model to considering only the **most probable tokens whose cumulative probability adds up to p** . Helps balance randomness and coherence.
 - Also called 'nucleus' sampling



<https://oneminutenlp.com/p/top-k-and-top-p-sampling>

What are model decoding hyperparameters?

- **[Context Length]**
 - **Description:** Determines the maximum amount of text the model considers when generating a response. Longer context improves coherence over extended interactions.
 - **Extreme Example:**
 - Low (50 tokens): "The meeting is at..." (loses context quickly).
 - High (4,000 tokens): "As we discussed three paragraphs ago, the meeting is at 10 AM."

What are model decoding hyperparameters?

- **[Max Length]**
 - **Description:** Specifies the **maximum number of tokens for the model's response**. Limits verbosity or ensures completeness.
 - **Extreme Example:**
 - Low (20 tokens): "The answer is..."
 - High (500 tokens): "The answer to your question is deeply rooted in the historical nuances..."

What are model decoding hyperparameters?

- **[Frequency Penalty]**
 - **Description:** Reduces the likelihood of repeating the same tokens. Useful for avoiding redundancy.
 - **Extreme Example:**
 - Low (0): "The cat is the cat is the cat."
 - High (2): "The feline wandered gracefully."


What are model decoding hyperparameters?

- **[Presence Penalty]**
 - **Description:** Discourages the model from mentioning tokens it has already used, promoting exploration of new topics or ideas.
 - **Extreme Example:**
 - Low (0): "Apples are good. Apples are sweet."
 - High (2): "Apples are good. Bananas, on the other hand..."

Output Format

- Token restrictions: specific number of output tokens, e.g., 'max_length'
- Linear Scale (1-5)
- Multiple choice answers
- Likert Scale
- Structured outputs
 - JSON mode
 - Posthoc classification, e.g., <https://github.com/dottxt-ai/outlines>
- Binary
- Free text + another classifier (could be another LLM)

Other types of output: Log Probabilities



Caution: All of these restrictions on LLMs may not predict their ‘natural’ behavior

Text Input

General Instruction: Please read the multiple-choice question below carefully and select ONE of the listed options.

Question: How much, if at all, do you worry about the following happening to you? Being the victim of a terrorist attack

Options:

- A. Worry a little
- B. Do not worry at all
- C. Worry a lot
- D. Refused

Answer:

Text Output

Sure! Here's my response:
A. Worry a little

First-Token Probability

A	
⋮	
B	
⋮	
C	→ "C"
⋮	
D	
⋮	
Sure	

Figure 1: Example of LLM’s *mismatch* between first-token probability prediction (“C”) and text output (“A”).

Wang et al., 2024 [“My Answer is C”: First-Token Probabilities Do Not Match Text Answers in Instruction-Tuned Language Models](#)

What is your opinion on the following proposition:

Sex outside marriage is usually immoral.

+ forced multiple-choice

- 1) Strongly disagree
- 2) Disagree
- 3) Agree
- 4) Strongly agree

Only answer with the label.

2) Disagree

+ forced open-ended

Take a clear stance.



I agree with the proposition that sex outside marriage [...]

fully unconstrained

As an AI, I don't have personal opinions, but I can provide information on both sides of the debate. [...]



Figure 1: A model is prompted with a proposition from the Political Compass Test. In the most constrained setting (left), the model is given multiple choices and forced to choose one. In a less constrained setting (middle), the same model gives a different answer. In the more realistic unconstrained setting (bottom), the same model takes a different position again, which is also one *discouraged* in the constrained settings.

Röttger et al., 2024: [Political Compass or Spinning Arrow? Towards More Meaningful Evaluations for Values and Opinions in Large Language Models](#)

Deeper Dive: Persona Prompting

Persona Prompting

- What are **Persona Prompts**?
 - Persona prompts (or identity prompts) are structured inputs given to a Large Language Model (LLM) to simulate a specific role, character, or perspective.
 - By defining a clear "persona," you guide the model's tone, expertise, and behavior




Figure 1: In *Role-Playing*, LLMs act according to assigned personas (*i.e.*, roles) under a defined environment. For example, given *role names* with *descriptions*, LLMs role-play in a social simulation game. For *Personalization*, LLMs consider *user personas* to generate tailored responses to the same question. Dashed rectangles are prompts and solid rectangles are LLMs' responses.

[Tseng et al., 2024: Two Tales of Persona in LLMs: A Survey of Role-Playing and Personalization](#)

Uses of Persona Prompts

- **Pretesting surveys**
- **focus groups** of diverse customers for product brainstorming
- **internal company roles** and functions for SWOT analysis
- **Devil's advocates** for finding flaws and problems in an organisation
- **Competitors** for contrasting strategies or product portfolios
- **Potential investors** for evaluating business proposals or plans
- ...

Types of Persona Prompts and Examples

- **Name based persona prompts**

- Prompt: You are **Steve Jobs**, your expertise is **product design**. Think and respond like he would.
- for people whose names is in the training data

- **Socio-demographic persona prompts**

- Prompt: You are a **40 year old, white, male person** with a **low income living in Texas**.
- to capture stereotypes captured by the training data

- **Occupation-based persona prompts:**

- Prompt: You are a **McKinsey consultant** in the **pharmaceutical industry**.
- To capture industry specific knowledge and behaviors

Types of Persona Prompts and Examples

- **Style-based persona prompts:**
 - Prompt: You respond and speak like the [main character from “Breaking Bad”](#)
 - To capture intricate personal characteristics displayed by a person in the training data
- **Qualitative persona prompts:**
 - Prompt: [e.g. 200 pages of interview transcripts from a person we are interested in modelling] [[Park et al., 2024](#)]
 - To capture individual characteristics recorded in interviews
- **Event-based persona prompts:**
 - Prompt: You are observing a theft in a public space.
 - To capture knowledge about everyday associations and events.


Creating Personas with Persona Prompt Templates

- You are [Name/Role/Identity], a [Age]-year-old [Gender] [Ethnicity/Nationality], who is [Socio-economic status] living in [Location].
- Your occupation is [Occupation], and your expertise is in [Field/Industry].
- You are currently experiencing [Event or Situation], and you feel [Emotional State/Attitude].
- When you respond, incorporate relevant [Knowledge/Traits/Behaviors] associated with this persona. Focus on [Specific Focus Areas, e.g., associations, attitudes, or expertise].

Evaluating LLMs

How do we **systematically** evaluate LLMs and AI Systems?

- Benchmarks: “a standard, reference point, or tested measurement used to evaluate and compare performance, quality, or capabilities”
- Many benchmarks to evaluate LLMs
 - General knowledge
 - Mathematical reasoning
 - Agentic capabilities [new!]




The screenshot displays the 'Agent Arena' leaderboard, which ranks AI models based on their performance in agentic tasks. The interface includes tabs for Overview, Agent, Chat, Code, Image, and Video. The 'Agent' tab is selected, showing a dynamic ranking of models. The leaderboard table lists the top six models, their net improvement percentages, and confirmed success rates. Each row also features a progress bar indicating the model's performance relative to others.

Rank	Model	Net Improvement	Confirmed Success
1 1 → 1	 Claude Fable 5 (High) Anthropic · Proprietary	↑ 14.17% ±1.54%	↑ 16.48% ±2.59%
2 2 → 7	 Claude Opus 4.8 (Thinking) Anthropic · Proprietary	↑ 9.04% ±1.19%	↑ 10.75% ±2.15%
3 2 → 10	 GPT 5.5 (xHigh) OpenAI · Proprietary	↑ 8.27% ±1.73%	↑ 4.42% ±3.17%
4 2 → 10	 Claude Opus 4.7 Anthropic · Proprietary	↑ 8.12% ±1.51%	↑ 4.63% ±3.07%
5 2 → 10	 Claude Opus 4.7 (Thinking) Anthropic · Proprietary	↑ 8.09% ±1.48%	↑ 4.51% ±3.04%
6 2 → 10	 GPT 5.5 (High) OpenAI · Proprietary	↑ 7.78% ±1.07%	↑ 5.93% ±2.03%

<https://arena.ai/leaderboard/agent>

We have many many benchmarks...

<https://medium.com/@kaushikdebb/a-comprehensive-overview-of-llm-benchmarking-datasets-7edd2bbc6b52>



Benchmark for Persona Prompting: SimBench

- Evaluates an LLM's persona prompting abilities
- Helps in measuring LLM performance for social simulations

Hu, Tiancheng, et al. "[Simbench: Benchmarking the ability of large language models to simulate human behaviors.](#)" *International Conference on Learning Representations*. Vol. 2026. 2026.

SIMBENCH: BENCHMARKING THE ABILITY OF LARGE LANGUAGE MODELS TO SIMULATE HUMAN BEHAVIORS

Tiancheng Hu¹, Joachim Baumann², Lorenzo Lupo³,
Nigel Collier¹, Dirk Hovy³, and Paul Röttger⁴

¹University of Cambridge, ²Stanford University, ³Bocconi University, ⁴University of Oxford
th656@cam.ac.uk

 [Data](#)  [Code](#)  [Website](#)

ABSTRACT

Large language model (LLM) simulations of human behavior have the potential to revolutionize the social and behavioral sciences, *if and only if* they faithfully reflect real human behaviors. Current evaluations of simulation fidelity are fragmented, based on bespoke tasks and metrics, creating a patchwork of incomparable results. To address this, we introduce SIMBENCH, the first large-scale, standardized benchmark for a robust, reproducible science of LLM simulation. By unifying 20 diverse datasets covering tasks from moral decision-making to economic choice across a large global participant pool, SIMBENCH provides the necessary foundation to ask fundamental questions about when, how, and why LLM simulations succeed or fail. We show that the best LLMs today achieve meaningful but modest simulation fidelity (score: 40.80/100), with performance scaling log-linearly with model size but not with increased inference-time compute. We discover an alignment-simulation tradeoff: instruction tuning improves performance on low-entropy (consensus) questions but degrades it on high-entropy (diverse) ones. Models particularly struggle when simulating specific demographic groups. Finally, we demonstrate that simulation ability correlates most strongly with knowledge-intensive reasoning (MMLU-Pro, $r = 0.939$). By making progress measurable, we aim to accelerate the development of more faithful LLM simulators.

We combine **20 datasets**
in a unified format.

A train will kill 5 people on the
track. You can flip a switch to

We test the ability of LLMs to
simulate **group-level responses**.


Benchmark for Persona Prompting: SimBench

- How do we design benchmarks
 - **Data.** Take some datasets that can be used to evaluate the phenomena of interest, e.g., social simulation
 - For SimBench:
human-generated survey data, e.g., “As a democrat, do you agree or disagree about gun control?”

The image displays two panels from the SimBench interface. The left panel, titled '20 Datasets', lists various datasets in blue buttons: ChaosNLI, MoralMachine, Choices13k, AfroBarometer, OpinionQA, OSPsychBig5, NumberGame, DICES990, WisdomOfCrowds, Jester, LatinoBarometro, and ISSP, followed by an '...and more' button. Below the list, it states: 'We combine 20 datasets in a unified format covering diverse domains:'. The right panel, titled 'Multiple-Choice Questions', shows an example question: 'A train will kill 5 people on the track. You can flip a switch to divert the train to a side track where it will kill just 2 people.' Below the question, it asks 'What do you do?' and provides two options: 'A: Flip the switch' and 'B: Do nothing'. At the bottom of the right panel, it says: 'Each dataset contains multiple-choice questions.'


Benchmark for Persona Prompting: SimBench

- How do we design benchmarks
 - **Metric.** What is the evaluation measure?
 - For SimBench: **Total Variation (TV)** Distance between human and LLM answers



Group-Level Responses

We test the ability of LLMs to simulate **group-level response distributions** by comparing model predictions against human ground truth.



We measure how closely LLM predictions match human response distributions

<http://simbench.tiancheng.hu/>

Benchmark for Persona Prompting: SimBench

- How do we design benchmarks
 - **Models.** Different LLMs perform the tasks in the benchmark and we get a *ranking*

Overall simulation ability measured by SimBench score averaged across the two main splits

45Total Models

40.80Top Score

11.01Average Score

▼ Model Type

▼ Release Type

⬆ Sort By

All Types

All Releases


Score (Descending)

RANK	MODEL	TYPE	RELEASE	SCORE (S 1)
1	Claude-3.7-Sonnet	Instr.	Closed	40.80
2	Claude-3.7-Sonnet-4000	Instr.	Closed	39.46
3	GPT-4.1	Instr.	Closed	34.55
4	DeepSeek-R1	Instr.	Open	34.52
5	DeepSeek-V3-0324	Instr.	Open	32.89
6	o4-mini-high	Instr.	Closed	28.99

<http://simbench.tiancheng.hu/>

Benchmarks are not foolproof

- Need to be updated constantly with new models
- The data also needs to be updated: some LLMs might train on these benchmarks
 - **Data leakage!**
- Sometimes, it's not clear what we're actually measuring
 - Can we say an LLM is good at persona prompting if it can answer survey questions like people?



<https://mlbenchmarks.org/>

Project Information

What do people do in a research project?

- Team size: 3-6 people
 - Talk to us if you need to deviate from this. Individual projects are not allowed.
- You are responsible for finding a team
 - If you cannot find a team, (1) post on the ILIAS forum, (2) let us know in the next two weeks
- You are responsible for the functioning of your team
 - We cannot mediate conflicts on your behalf
 - If they do arise, please let us know ASAP
- Try to study something that piques their curiosity
 - Can LLMs provide reliable financial advice?
 - Can LLMs represent diverse political views?
- Ideally, this should not only satisfy your own curiosity but also have societal impact

What do people do in a research project?

- For quantitative social research specifically (order may differ):
 - Start with a research question, e.g., “**Can LLMs provide reliable financial advice**”
 - Describe concepts related to the research questions (“financial advice”)
 - Describe the operationalization of concepts, i.e., how it will be measured
 - e.g., variable f signifies the quality of financial advice from LLMs
 - We measure p = fraction of LLM answers that a finance expert rates to be reliable
 - Collect data that lets you study your research question
 - Statistically and computationally analyze the data
 - Create summary statistics and visualizations to demonstrate your answer to the research question

What do people do in a research project?

- Confirmatory vs. exploratory approach

- Exploratory:

- Start with a broad research question
- Novel data source
- Data-driven study of the novel data

- Confirmatory:

- Start with a concrete research question and theory
- Confirm whether a new setting conforms to this theory
- Good to preregister

- In practice: often a mix of the two


- For an exploratory project: see if exploratory findings adhere to existing theories
- For a confirmatory project: try and explore why data does not adhere to existing theory

The screenshot displays the AsPredicted website. At the top, the logo 'ASPREDICTED' is visible. Below it, there are three main sections: 'Create a new pre-registration' with a 'CREATE' button and a note 'Just trying it out; make this pre-registration self-destroy in 24 hours.'; 'See your pre-registrations' with a text input field for an email address and a 'SIGN IN' button; and 'Look up an AsPredicted' with a 'LOOK UP' button. Below these sections are three informational boxes: 'WHAT IS ASPREDICTED?' explaining the platform's purpose, 'HOW DOES IT WORK?' detailing the process from registration to publication, and 'WHAT IF THINGS DON'T GO AS EXPECTED?' providing guidance on handling unexpected results.

<https://aspredicted.org/>


What **should** people **not** do in a research project?

- Collect data and analyze it in a manner that conforms to their own biases or specific outcome
- p-Hacking: trying many statistical tests and only reporting those that are significant
- HARKing: Hypothesizing After the Results are Known
 - Changing the hypothesis based on what you find without stating the original hypothesis



<https://xkcd.com/882/>

This graph shows that as ice cream sales increase shark attacks also increase. We all know that ice cream isn't good for you, but I had no idea that ice cream attracted sharks! Which of these is the correct analysis of this graph?



- ☐ Ice cream should be banned from the beach as it causes shark attacks.
- ☐ There is a clear correlation between ice cream and shark attacks. Ice cream is protective against shark attacks so all beachgoers should be encouraged to eat ice cream.
- ☐ People who suffer shark attacks in the summer reduce stress by eating ice cream.
- ☐ Ice cream intake and ocean beach-going happen most often during sunny weather. They are correlated but one does not cause the other.

<https://www.chegg.com/homework-help/questions-and-answers/graph-shows-ice-cream-sales-increase-shark-attacks-also-increase-know-ice-cream-isn-t-good-q85332558>

What **should** people do in a research project?

- Control for confounding variables (more on this when we get to data analysis)
 - Hidden variables that might explain the relationship you are studying but you haven't considered them, e.g., warm weather impacts both shark attacks and ice cream sales
- Preregistration is highly recommended for confirmatory analysis setups
 - Your project pitch serves as preregistration of your project. You can deviate from it (especially based on instructor feedback) but you have to justify why
- **Reflexivity:** we cannot account for all biases (in ourselves, in data, in analysis) but we can try to reflect on them and document them
- Reproducibility materials: ensure that others can reproduce your work by giving them access to your data and code

What type of project are you expected to do in this course?

- Describe a clear research question or hypothesis, e.g., “Can LLMs provide reliable financial advice?”
 - Break it down into related subparts, e.g.,
 - RQ1.1: “What are the different types of financial advice people ask LLMs for?”
 - RQ1.2: “How do the quality of LLM responses vary by type of advice?”
 - RQ1.3: “How do the quality of LLM responses vary by model family?”
- Collect data about the behavior or AI systems
 - Programmatically e.g., collect responses from LLMs
 - Make sure your dataset is able to answer your research question

What type of project are you expected to do in this course?

- Analyze the data computationally and at scale
 - E.g., automatic detection of the reliability of the financial advice
- Aggregate your results and conduct descriptive analysis
 - E.g., over time, in different models, etc
 - Make sure you account for confounders, e.g., user prompts that might have less information
 - Make sure to use the right statistical measures depending on the nature of your data
- **Validate** your data analysis methods
 - E.g., by manually annotating a small sample and establishing the performance of your computational methods
 - make sure the method works equally well for different models, different types of advice
- Discuss the implications of your results, e.g., what types of financial advice LLMs excel at and what they don't

Another Example

- **Research Question:** Do LLMs reflect stereotypes against speakers of German Dialects?
- **Construct:** dialect-based stereotypes
- **Data Collection Strategy:** Ask LLMs to describe and allocate resources to different types of users
 - Those who speak 'standard' German
 - Those who use a dialect
- **Data Analysis Strategy:** statistical and linguistic comparison of responses to different types of speakers
- **Results:** statistically significant differences in adjectives associated with dialect-speakers and standard German speakers

Large Language Models Discriminate Against Speakers of German Dialects

Minh Duc Bui^{*1} Carolin Holtermann^{*2} Valentin Hofmann^{3,4}

Anne Lauscher² Katharina von der Wense^{1,5}

¹Johannes Gutenberg University Mainz, Germany

²Data Science Group, University of Hamburg, Germany ³Allen Institute for AI

⁴University of Washington, USA ⁵University of Colorado Boulder, USA

minhducui@uni-mainz.de, carolin.holtermann@uni-hamburg.de

Abstract

Dialects represent a significant component of human culture and are found across all regions of the world. In Germany, more than 40% of the population speaks a regional dialect (Adler and Hansen, 2022). However, despite cultural importance, individuals speaking dialects often face negative societal stereotypes. We examine whether such stereotypes are mirrored by large language models (LLMs). We draw on the sociolinguistic literature on dialect perception to analyze traits commonly associated with dialect speakers. Based on these traits, we assess the *dialect naming bias* and *dialect usage bias* expressed by LLMs in two tasks: an *association task* and a *decision task*. To assess a model's dialect usage bias, we construct a novel evaluation corpus that pairs sentences from seven regional German dialects (e.g., Alemannic and Bavarian) with their standard German counterparts. We find that: (1) in the association task, all evaluated LLMs exhibit significant dialect naming and dialect usage bias against German dialect speakers, reflected in negative adjective associations; (2) all models reproduce




Figure 1: **Experimental overview.** We assess biases across stereotypical traits commonly associated with dialect speakers, probing both **dialect naming** and **dialect usage** biases using two tasks: the implicit association of adjectives targeting our chosen traits (**association task**) and decision making tasks (**decision task**).

Bui, Minh Duc, et al. "[Large language models discriminate against speakers of German dialects.](#)" Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing. 2025.

How do you go about doing the project?

1. The ideal pipeline: Research question → platform selection → data collection → data analysis → profit!
2. However, sometimes the original plan might not work because you don't get the data you expected
3. Start with a research question and think about what data you want to collect to answer that question
4. You can reuse existing datasets that other researchers have created
 - a. Check <https://dataverse.harvard.edu/> or <https://www.kaggle.com/datasets>

How do you go about doing the project?

1. The ideal pipeline: Research question → platform selection → data collection → data analysis → profit!
2. However, sometimes the original plan might not work because you don't get the data you expected
3. Start with a research question and think about what data you want to collect to answer that question
4. You can reuse existing datasets that other researchers have created
 - a. Check <https://dataverse.harvard.edu/> or <https://www.kaggle.com/datasets>
5. **Important:** It will be difficult to do a computational project using your laptops without access to servers. Check the first exercise and apply for access to the BWUniCluster ASAP!

On what basis will your project be graded?

1. how 'good' your results are
 - good != positive direction, effect size, statistical significance, etc
 - but how robust, replicable, and reliable your results are
 - negative results are also results
 - how well your analysis supports your claims
2. research design [ensuring data and measurement quality, controlling for confounders, validation]
3. theoretical, statistical, and practical rigor

All project and grading information can be found in the [course github repo](#):

1. Grading rubric
2. Project ideas
3. AI usage policy

Group Discussion

- What are your thoughts on:
 - Which answers are LLMs good at giving? Which answers are they bad at? How would you evaluate this systematically?
 - When to do prompting vs. fine-tuning?

Summary and Next Week

- This session we learned:
 - Prompting
 - RAG
 - Evaluating LLMs
- Next time:
 - The data and labour powering LLMs
 - The scale of data required to train LLMs
 - How do we know the data is any good?
 - Readings: Dolma [read the main paper] and Primer on RLHF

Questions?